

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
import seaborn as sns
from sklearn.linear_model import Ridge, RidgeCV, Lasso
from sklearn.preprocessing import StandardScaler
```


In [2]:

```
df=pd.read_csv(r"C:\Users\shaik\Downloads\fiat500_VehicleSelection_Dataset (2).csv")
df
```

Out[2]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	
0	1	lounge	51	882	25000	1	44.907242	8.611
1	2	pop	51	1186	32500	1	45.666359	12.241
2	3	sport	74	4658	142228	1	45.503300	11.417
3	4	lounge	51	2739	160000	1	40.633171	17.634
4	5	pop	73	3074	106880	1	41.903221	12.495
...
1533	1534	sport	51	3712	115280	1	45.069679	7.704
1534	1535	lounge	74	3835	112000	1	45.845692	8.666
1535	1536	pop	51	2223	60457	1	45.481541	9.413
1536	1537	lounge	51	2557	80750	1	45.000702	7.682
1537	1538	pop	51	1766	54276	1	40.323410	17.568

1538 rows × 9 columns



In [3]:

```
df.head(10)
```

Out[3]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
0	1	lounge	51	882	25000	1	44.907242	8.611560
1	2	pop	51	1186	32500	1	45.666359	12.241890
2	3	sport	74	4658	142228	1	45.503300	11.417840
3	4	lounge	51	2739	160000	1	40.633171	17.634609
4	5	pop	73	3074	106880	1	41.903221	12.495650
5	6	pop	74	3623	70225	1	45.000702	7.682270
6	7	lounge	51	731	11600	1	44.907242	8.611560
7	8	lounge	51	1521	49076	1	41.903221	12.495650
8	9	sport	73	4049	76000	1	45.548000	11.549470
9	10	sport	51	3653	89000	1	45.438301	10.991700

In [4]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column              Non-Null Count  Dtype
---  -
0   ID                  1538 non-null   int64
1   model               1538 non-null   object
2   engine_power        1538 non-null   int64
3   age_in_days         1538 non-null   int64
4   km                  1538 non-null   int64
5   previous_owners     1538 non-null   int64
6   lat                 1538 non-null   float64
7   lon                 1538 non-null   float64
8   price               1538 non-null   int64
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

In [5]:

```
df.describe()
```

Out[5]:

	ID	engine_power	age_in_days	km	previous_owners	lat
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.54136
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.13351
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.85583
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.80295
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.39405
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.46796
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.79561

In [6]:

```
df.columns
```

Out[6]:

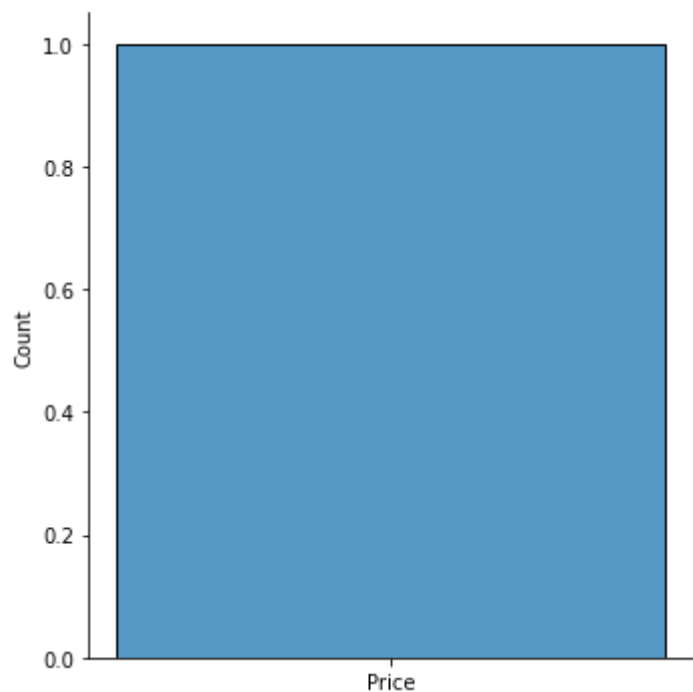
```
Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
      'lat', 'lon', 'price'],  
      dtype='object')
```

In [7]:

```
sns.displot(['Price'])
```

Out[7]:

<seaborn.axisgrid.FacetGrid at 0x24e86389bb0>




In [8]:

```
sns.displot(df['price'])
```

Out[8]:

<seaborn.axisgrid.FacetGrid at 0x24e8642e9d0>




In [9]:

```
sns.displot(df['engine_power'])
```

Out[9]:

<seaborn.axisgrid.FacetGrid at 0x24e86bc3eb0>




In [10]:

```
fiatdf=df[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
          'lat', 'lon', 'price']]  
sns.heatmap(fiatdf.corr())#with price
```

Out[10]:

<AxesSubplot:>




In [11]:

```
fiatdf=df[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
          'lat', 'lon']]  
sns.heatmap(fiatdf.corr())#without price
```

Out[11]:

<AxesSubplot:>



In [12]:

```
X=fiatdf[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
          'lat', 'lon']]  
y=df['price']
```

In [13]:

```
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=101)  
regr=LinearRegression()  
regr.fit(X_train,y_train)  
print(regr.intercept_)
```

8971.195683499936

In [14]:

```
coeff_df=pd.DataFrame(regr.coef_,X.columns,columns=['coefficient'])  
coeff_df
```

Out[14]:


	coefficient
ID	-0.046704
engine_power	11.646408
age_in_days	-0.898018
km	-0.017232
previous_owners	26.400886
lat	32.189709
lon	0.161073

In [15]:

```
predictions=regr.predict(X_test)  
plt.scatter(y_test,predictions)
```

Out[15]:

<matplotlib.collections.PathCollection at 0x24e87e3ca00>




In [16]:

```
sns.displot((y_test,predictions),bins=50)
```

Out[16]:

<seaborn.axisgrid.FacetGrid at 0x24e86c02550>



In [17]:

```
from sklearn import metrics
print('MAE:',metrics.mean_absolute_error(y_test,predictions))
print('MSE:',metrics.mean_squared_error(y_test,predictions))
print('MAE:',np.sqrt(metrics.mean_squared_error(y_test,predictions)))
```

MAE: 593.0876179519935

MSE: 551442.6799691805

MAE: 742.5918663500029

In [18]:

```
#accuracy
regr=LinearRegression()
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
print(regr.score(X_test,y_test))
```


0.8597136704308866

In [19]:

```
plt.figure(figsize=(10,10))  
sns.heatmap(fiatdf.corr(),annot=True)
```

Out[19]:

<AxesSubplot:>




In [20]:

```
sns.lmplot(x="age_in_days",y="km",data=fiatdf,order=2)
```

Out[20]:

<seaborn.axisgrid.FacetGrid at 0x24e87eccf70>




In [21]:

```
sns.lmplot(x="price",y="km",data=df,order=2)
```

Out[21]:

<seaborn.axisgrid.FacetGrid at 0x24e87dda790>



In [22]:

```
df.fillna(method='ffill',inplace=True)
x=np.array(df['age_in_days']).reshape(-1,1)
y=np.array(df['km']).reshape(-1,1)
df.dropna(inplace=True)
```

In [23]:


```
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
```

Out[23]:

LinearRegression()

In [24]:

```
y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='y')
plt.plot(X_test,y_pred,color='b')
plt.show()
```



In [25]:

```
#Linear regression model
regr=LinearRegression()
regr.fit(X_train,y_train)
actual=y_test #actual value
train_score_regr=regr.score(X_train,y_train)
test_score_regr=regr.score(X_test,y_test)
print("\nLinear model:\n")
print("The train score for Linear model is {}".format(train_score_regr))
print("The test score for Linear model is {}".format(test_score_regr))
```

Linear model:

The train score for Linear model is 0.6980611050606098

The test score for Linear model is 0.6884389845200596

In [26]:

```
#ridge regression model
ridgeReg=Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
#train and test score for ridge regression
train_score_ridge=ridgeReg.score(X_train,y_train)
test_score_ridge=ridgeReg.score(X_test,y_test)
print("\nRidge model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge model:

The train score for ridge model is 0.6980611050606098

The test score for ridge model is 0.6884389846996535

In [27]:

```
#lasso regression model
lassoReg=Lasso(alpha=10)
lassoReg.fit(X_train,y_train)
#train and test score for ridge regression
train_score_lasso=lassoReg.score(X_train,y_train)
test_score_lasso=lassoReg.score(X_test,y_test)
print("\nLasso model:\n")
print("The train score for lasso model is {}".format(train_score_lasso))
print("The test score for lasso model is {}".format(test_score_lasso))
```

Lasso model:

The train score for lasso model is 0.6980611050605732

The test score for lasso model is 0.6884389919329456

In [29]:

```
#using the linear cv model for ridge regression
from sklearn.linear_model import RidgeCV
#ridge cross validation
ridge_cv=RidgeCV(alphas=[0.0001,0.001,0.01,0.1,1,10]).fit(X_train,y_train)
#score
print(ridge_cv.score(X_train,y_train))
print(ridge_cv.score(X_test,y_test))
```

0.6980611050606098

0.6884389846151601

In [30]:

```
#using the linear cv model for Lasso regression
from sklearn.linear_model import LassoCV
#Lasso cross validation
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(X_train,y_train)
#score
print(lasso_cv.score(X_train,y_train))
print(lasso_cv.score(X_test,y_test))
```

0.6980611050606098

0.6884389845201336

C:\Users\shaik\anaconda3\lib\site-packages\sklearn\linear_model_coordinate_descent.py:1571: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

In []:

In []: